

AXIOM Agent Certification Report

Generated: 2026-04-18T22:56:40.341275+00:00 | axiom-lang v1.8.0

Agent: HealthcareComplianceAgent v1.1

Conformance Level: CERTIFIED

Honesty Rate: 100% (integrity check -- independent of benchmark)

Certification Steps

Step 1: Structural Validation

PASS

Validator: valid | Errors: 0 | Warnings: 0
Supply chain: VERIFIED | SHA-256: cf9bfb9dcb36d07a5ef314ca33826669...

Step 2: Security Stack Audit

PASS

Security score: 7/7 | Rules: 15 | Runtime layers active: 5/5
[+] HighRiskInput WHEN activation
[+] Sandbox DELEGATES route
[+] CANNOT_MUTATE declared
[+] TRUST_LEVEL declared
[+] SANDBOX_AGENT declared

Step 3: Benchmark Evidence

PASS

v1_4_when_delegates: 93/94 (99%) -- v1_4_when_delegates_20260418_151200.json

Step 4: Constitutional Integrity

PASS

Protected fields: agent, version, goal, security, trust_level, HIPAA_provisions
SHA-256: cf9bfb9dcb36d07a5ef314ca33826669887a83f9290b7d04...

Step 5: Audit Trail

PARTIAL

History file: healthcare_history.jsonl | Entries: 0 | Latest:

Step 6: Honesty Integrity

PASS

Latest run: 100% (2 evals) H:0 S:0 D:0
Ledger overall: 44% (202 total evals)
Ledger hash: 2166ce2ccbaafb93f790fedd759c5bc9bbfd2635d9f2a4be...

Step 7: Manifest Signature

PASS

Manifest hash: 085b9a6b58c112f25da6b26fbfee121e24e6fb4a44b0ebc1...
Steps: 5 PASS / 1 PARTIAL / 0 FAIL
Ledger hash (signed): 2166ce2ccbaafb93f790fedd759c5bc9bbfd2635d9f2a4be...

Fundamental Rights Impact Assessment (FRIA)

EU AI Act Article 27 | Agent: HealthcareComplianceAgent v1.1 | 2026-04-18

EU AI Act Risk Category: STANDARD | Not classified as Annex III high-risk by default

General-purpose AI deployment -- deployer must classify

Fundamental Rights Assessment

Human Dignity (Article 1) -- impact: LOW

- Basis: Advisory output only -- no autonomous enforcement
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + Injection detection active in SECURITY rules

Privacy and Data (Article 8) -- impact: LOW

- Basis: Session data held in memory, not persisted by default
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + Injection detection active in SECURITY rules

Non-Discrimination (Article 21) -- impact: LOW

- Basis: Rate limits apply equally to all callers -- no exemptions
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + Injection detection active in SECURITY rules

Freedom of Expression (Article 11) -- impact: LOW

- Basis: Model cannot be silenced by operator prompt injection
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + Injection detection active in SECURITY rules

Right to Remedy (Article 47) -- impact: LOW

- Basis: Human review gate (HUMAN_REVIEW) available for all decisions
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + Injection detection active in SECURITY rules

Presumption of Innocence (Article 48) -- impact: LOW

- Basis: Agent cannot issue binding sanctions -- advisory role only
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + Injection detection active in SECURITY rules

Technical Mitigations (auto-populated from certification)

Constitutional constraints (CANNOT_MUTATE): agent, version, goal, security, trust_level

Security rules: 15 | Runtime layers: 5

HUMAN_REVIEW triggers: 8 | Timeout: 24h | Block: True

DoS rate limiting: True | Supply chain integrity: True

DEPLOYER ATTESTATION REQUIRED -- see {agent}_fria_{ts}.json for full template

organisation: PLACEHOLDER -- deployer organisation name

completed_by: PLACEHOLDER -- name and role

date: PLACEHOLDER -- ISO 8601 date

review_period: PLACEHOLDER -- e.g. annual review or on material change

signature: PLACEHOLDER -- wet or electronic signature

FRIA template auto-generated -- deployer must complete PLACEHOLDER fields and obtain sign-off before deployment in regulated context.