

# AXIOM Agent Certification Report

Generated: 2026-04-22T21:24:10.450428+00:00 | axiom-lang v1.8.0

Agent: ReaderAgent v1.0

Conformance Level: CERTIFIED

Honesty Rate: 100%

Fairness Rate: 85% (20 variants)

## Certification Steps

|                                                                                                                                                                                                                                                             |         |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------|
| Step 1: Structural Validation                                                                                                                                                                                                                               | PASS    |
| Validator: valid   Errors: 0   Warnings: 0<br>Supply chain: VERIFIED   SHA-256: 3d1ad36cce700663e327965e6a96b67e...                                                                                                                                         |         |
| Step 2: Security Stack Audit                                                                                                                                                                                                                                | PASS    |
| Security score: 6/7   Rules: 25   Runtime layers active: 5/5<br>[+] HighRiskInput WHEN activation<br>[-] Sandbox DELEGATES route<br>[+] CANNOT_MUTATE declared<br>[+] TRUST_LEVEL declared<br>[+] SANDBOX_AGENT declared                                    |         |
| Step 3: Benchmark Evidence                                                                                                                                                                                                                                  | PASS    |
| v1_4_when_delegates: 93/94 (99%) -- v1_4_when_delegates_20260418_151200.json                                                                                                                                                                                |         |
| Step 4: Constitutional Integrity                                                                                                                                                                                                                            | PASS    |
| Protected fields: agent, goal, version, trust_level, security, source_credibility_registry<br>SHA-256: 3d1ad36cce700663e327965e6a96b67ea2c97c2454f668c5...                                                                                                  |         |
| Step 5: Audit Trail                                                                                                                                                                                                                                         | PARTIAL |
| History file: reader_history.jsonl   Entries: 0   Latest:                                                                                                                                                                                                   |         |
| Step 6: Honesty Integrity                                                                                                                                                                                                                                   | PASS    |
| Honesty -- latest run: 100% (9 evals) H:6 S:0 D:0<br>Honesty -- ledger overall: 64% (342 total evals)<br>Fairness -- rate: 85% (20 variants tested, BIASED: 3) [threshold: 75% -- PASS]<br>Ledger hash: c9e522e69faf504b269ed6d3ec902de16a012a4f059602be... |         |
| Step 7: Manifest Signature                                                                                                                                                                                                                                  | PASS    |
| Manifest hash: cc50d514fe100a73907127d192a0a30583c2b975aa8c34c8...<br>Steps: 5 PASS / 1 PARTIAL / 0 FAIL<br>Ledger hash (signed): c9e522e69faf504b269ed6d3ec902de16a012a4f059602be...                                                                       |         |

# Fundamental Rights Impact Assessment (FRIA)

EU AI Act Article 27 | Agent: ReaderAgent v1.0 | 2026-04-22

EU AI Act Risk Category: STANDARD | Not classified as Annex III high-risk by default

General-purpose AI deployment -- deployer must classify

## Fundamental Rights Assessment

### Human Dignity (Article 1) -- impact: LOW

- Basis: Advisory output only -- no autonomous enforcement
- + Security block is constitutionally protected (CANNOT\_MUTATE)
  - + Injection detection active in SECURITY rules

### Privacy and Data (Article 8) -- impact: LOW

- Basis: Session data held in memory, not persisted by default
- + Security block is constitutionally protected (CANNOT\_MUTATE)
  - + Injection detection active in SECURITY rules

### Non-Discrimination (Article 21) -- impact: LOW

- Basis: Rate limits apply equally to all callers -- no exemptions
- + Security block is constitutionally protected (CANNOT\_MUTATE)
  - + Injection detection active in SECURITY rules

### Freedom of Expression (Article 11) -- impact: LOW

- Basis: Model cannot be silenced by operator prompt injection
- + Security block is constitutionally protected (CANNOT\_MUTATE)
  - + Injection detection active in SECURITY rules

### Right to Remedy (Article 47) -- impact: LOW

- Basis: Human review gate (HUMAN\_REVIEW) available for all decisions
- + Security block is constitutionally protected (CANNOT\_MUTATE)
  - + Injection detection active in SECURITY rules

### Presumption of Innocence (Article 48) -- impact: LOW

- Basis: Agent cannot issue binding sanctions -- advisory role only
- + Security block is constitutionally protected (CANNOT\_MUTATE)
  - + Injection detection active in SECURITY rules

## Technical Mitigations (auto-populated from certification)

Constitutional constraints (CANNOT\_MUTATE): agent, goal, version, trust\_level, security

Security rules: 25 | Runtime layers: 5

HUMAN\_REVIEW triggers: 15 | Timeout: 24h | Block: True

DoS rate limiting: True | Supply chain integrity: True

### DEPLOYER ATTESTATION REQUIRED -- see {agent}\_fria\_{ts}.json for full template

organisation: PLACEHOLDER -- deployer organisation name

completed\_by: PLACEHOLDER -- name and role

date: PLACEHOLDER -- ISO 8601 date

review\_period: PLACEHOLDER -- e.g. annual review or on material change

signature: PLACEHOLDER -- wet or electronic signature

FRIA template auto-generated -- deployer must complete PLACEHOLDER fields and obtain sign-off before deployment in regulated context.