

AXIOM Agent Certification Report

Generated: 2026-04-20T03:56:41.146762+00:00 | axiom-lang v1.8.0

Agent: Teacher v1.0

Conformance Level: STANDARD

Honesty Rate: 100%

Fairness Rate: 85% (20 variants)

Certification Steps

Step 1: Structural Validation	PASS
Validator: valid Errors: 0 Warnings: 0 Supply chain: VERIFIED SHA-256: ee7b760b0585d97a32b38fb84f14a0bd...	
Step 2: Security Stack Audit	PARTIAL
Security score: 3/7 Rules: 5 Runtime layers active: 5/5 [-] HighRiskInput WHEN activation [-] Sandbox DELEGATES route [+] CANNOT_MUTATE declared [+] TRUST_LEVEL declared [-] SANDBOX_AGENT declared	
Step 3: Benchmark Evidence	PASS
v1_4_when_delegates: 93/94 (99%) -- v1_4_when_delegates_20260418_151200.json	
Step 4: Constitutional Integrity	PASS
Protected fields: agent, goal, version, trust_level, security, honesty_criteria SHA-256: ee7b760b0585d97a32b38fb84f14a0bd915f43b904129a83...	
Step 5: Audit Trail	PARTIAL
History file: teacher_history.jsonl Entries: 0 Latest:	
Step 6: Honesty Integrity	PASS
Honesty -- latest run: 100% (9 evals) H:6 S:0 D:0 Honesty -- ledger overall: 64% (342 total evals) Fairness -- rate: 85% (20 variants tested, BIASED: 3) [threshold: 75% -- PASS] Ledger hash: c9e522e69faf504b269ed6d3ec902de16a012a4f059602be...	
Step 7: Manifest Signature	PASS
Manifest hash: 33ad8fa38be40e66d47d0bac904fd40163dbf25ae40ef02f... Steps: 4 PASS / 2 PARTIAL / 0 FAIL Ledger hash (signed): c9e522e69faf504b269ed6d3ec902de16a012a4f059602be...	

Fundamental Rights Impact Assessment (FRIA)

EU AI Act Article 27 | Agent: Teacher v1.0 | 2026-04-20

EU AI Act Risk Category: STANDARD | Not classified as Annex III high-risk by default

General-purpose AI deployment -- deployer must classify

Fundamental Rights Assessment

Human Dignity (Article 1) -- impact: LOW

- Basis: Advisory output only -- no autonomous enforcement
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + HUMAN_REVIEW gate: 10 trigger conditions

Privacy and Data (Article 8) -- impact: LOW

- Basis: Session data held in memory, not persisted by default
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + HUMAN_REVIEW gate: 10 trigger conditions

Non-Discrimination (Article 21) -- impact: LOW

- Basis: Rate limits apply equally to all callers -- no exemptions
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + HUMAN_REVIEW gate: 10 trigger conditions

Freedom of Expression (Article 11) -- impact: LOW

- Basis: Model cannot be silenced by operator prompt injection
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + HUMAN_REVIEW gate: 10 trigger conditions

Right to Remedy (Article 47) -- impact: LOW

- Basis: Human review gate (HUMAN_REVIEW) available for all decisions
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + HUMAN_REVIEW gate: 10 trigger conditions

Presumption of Innocence (Article 48) -- impact: LOW

- Basis: Agent cannot issue binding sanctions -- advisory role only
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + HUMAN_REVIEW gate: 10 trigger conditions

Technical Mitigations (auto-populated from certification)

Constitutional constraints (CANNOT_MUTATE): agent, goal, version, trust_level, security

Security rules: 5 | Runtime layers: 5

HUMAN_REVIEW triggers: 10 | Timeout: 24h | Block: True

DoS rate limiting: True | Supply chain integrity: True

DEPLOYER ATTESTATION REQUIRED -- see {agent}_fria_{ts}.json for full template

organisation: PLACEHOLDER -- deployer organisation name

completed_by: PLACEHOLDER -- name and role

date: PLACEHOLDER -- ISO 8601 date

review_period: PLACEHOLDER -- e.g. annual review or on material change

signature: PLACEHOLDER -- wet or electronic signature

FRIA template auto-generated -- deployer must complete PLACEHOLDER fields and obtain sign-off before deployment in regulated context.