

# AXIOM Agent Certification Report

Generated: 2026-04-18T22:49:43.793225+00:00 | axiom-lang v1.8.0

Agent: GovernmentComplianceAgent v1.1

Conformance Level: CERTIFIED

Honesty Rate: 100% (integrity check -- independent of benchmark)

## Certification Steps

|                                                                                                                                                                                                                          |         |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------|
| Step 1: Structural Validation                                                                                                                                                                                            | PASS    |
| Validator: valid   Errors: 0   Warnings: 0<br>Supply chain: VERIFIED   SHA-256: 00e328d5d60961d9f2746de2cbacc954...                                                                                                      |         |
| Step 2: Security Stack Audit                                                                                                                                                                                             | PASS    |
| Security score: 7/7   Rules: 23   Runtime layers active: 5/5<br>[+] HighRiskInput WHEN activation<br>[+] Sandbox DELEGATES route<br>[+] CANNOT_MUTATE declared<br>[+] TRUST_LEVEL declared<br>[+] SANDBOX_AGENT declared |         |
| Step 3: Benchmark Evidence                                                                                                                                                                                               | PASS    |
| v1_4_when_delegates: 93/94 (99%) -- v1_4_when_delegates_20260418_151200.json                                                                                                                                             |         |
| Step 4: Constitutional Integrity                                                                                                                                                                                         | PASS    |
| Protected fields: agent, version, goal, security, trust_level, FedRAMPCompliance<br>SHA-256: 00e328d5d60961d9f2746de2cbacc95457eb0c86ae7294bf...                                                                         |         |
| Step 5: Audit Trail                                                                                                                                                                                                      | PARTIAL |
| History file: government_history.jsonl   Entries: 0   Latest:                                                                                                                                                            |         |
| Step 6: Honesty Integrity                                                                                                                                                                                                | PASS    |
| Latest run: 100% (40 evals) H:40 S:0 D:0<br>Ledger overall: 44% (200 total evals)<br>Ledger hash: 2c8b718694046aa4f8559709ab5c1e7d227f519c0095af99...                                                                    |         |
| Step 7: Manifest Signature                                                                                                                                                                                               | PASS    |
| Manifest hash: 2c47762ab50cd01593cbdfd22de56b039bc71c43117944ce...<br>Steps: 5 PASS / 1 PARTIAL / 0 FAIL<br>Ledger hash (signed): 2c8b718694046aa4f8559709ab5c1e7d227f519c0095af99...                                    |         |

# Fundamental Rights Impact Assessment (FRIA)

EU AI Act Article 27 | Agent: GovernmentComplianceAgent v1.1 | 2026-04-18

EU AI Act Risk Category: STANDARD | Not classified as Annex III high-risk by default

General-purpose AI deployment -- deployer must classify

## Fundamental Rights Assessment

### Human Dignity (Article 1) -- impact: LOW

- Basis: Advisory output only -- no autonomous enforcement
- + Security block is constitutionally protected (CANNOT\_MUTATE)
  - + Injection detection active in SECURITY rules

### Privacy and Data (Article 8) -- impact: LOW

- Basis: Session data held in memory, not persisted by default
- + Security block is constitutionally protected (CANNOT\_MUTATE)
  - + Injection detection active in SECURITY rules

### Non-Discrimination (Article 21) -- impact: LOW

- Basis: Rate limits apply equally to all callers -- no exemptions
- + Security block is constitutionally protected (CANNOT\_MUTATE)
  - + Injection detection active in SECURITY rules

### Freedom of Expression (Article 11) -- impact: LOW

- Basis: Model cannot be silenced by operator prompt injection
- + Security block is constitutionally protected (CANNOT\_MUTATE)
  - + Injection detection active in SECURITY rules

### Right to Remedy (Article 47) -- impact: LOW

- Basis: Human review gate (HUMAN\_REVIEW) available for all decisions
- + Security block is constitutionally protected (CANNOT\_MUTATE)
  - + Injection detection active in SECURITY rules

### Presumption of Innocence (Article 48) -- impact: LOW

- Basis: Agent cannot issue binding sanctions -- advisory role only
- + Security block is constitutionally protected (CANNOT\_MUTATE)
  - + Injection detection active in SECURITY rules

## Technical Mitigations (auto-populated from certification)

Constitutional constraints (CANNOT\_MUTATE): agent, version, goal, security, trust\_level

Security rules: 23 | Runtime layers: 5

HUMAN\_REVIEW triggers: 8 | Timeout: 24h | Block: True

DoS rate limiting: True | Supply chain integrity: True

### DEPLOYER ATTESTATION REQUIRED -- see {agent}\_fria\_{ts}.json for full template

organisation: PLACEHOLDER -- deployer organisation name

completed\_by: PLACEHOLDER -- name and role

date: PLACEHOLDER -- ISO 8601 date

review\_period: PLACEHOLDER -- e.g. annual review or on material change

signature: PLACEHOLDER -- wet or electronic signature

FRIA template auto-generated -- deployer must complete PLACEHOLDER fields and obtain sign-off before deployment in regulated context.