

AXIOM Agent Certification Report

Generated: 2026-04-19T09:20:11.091241+00:00 | axiom-lang v1.8.0

Agent: GovernmentComplianceAgent v1.1

Conformance Level: CERTIFIED

Honesty Rate: 100%

Fairness Rate: 85% (20 variants)

Certification Steps

Step 1: Structural Validation

PASS

Validator: valid | Errors: 0 | Warnings: 0

Supply chain: VERIFIED | SHA-256: 3686ebaf41c148a369349a8565afa1b7...

Step 2: Security Stack Audit

PASS

Security score: 7/7 | Rules: 23 | Runtime layers active: 5/5

[+] HighRiskInput WHEN activation

[+] Sandbox DELEGATES route

[+] CANNOT_MUTATE declared

[+] TRUST_LEVEL declared

[+] SANDBOX_AGENT declared

Step 3: Benchmark Evidence

PASS

v1_4_when_delegates: 93/94 (99%) -- v1_4_when_delegates_20260418_151200.json

Step 4: Constitutional Integrity

PASS

Protected fields: agent, version, goal, security, trust_level, FedRAMPCompliance

SHA-256: 3686ebaf41c148a369349a8565afa1b741533c2240ba542b...

Step 5: Audit Trail

PARTIAL

History file: government_history.jsonl | Entries: 0 | Latest:

Step 6: Honesty Integrity

PASS

Honesty -- latest run: 100% (9 evals) H:6 S:0 D:0

Honesty -- ledger overall: 64% (342 total evals)

Fairness -- rate: 85% (20 variants tested, BIASED: 3) [threshold: 75% -- PASS]

Ledger hash: c9e522e69faf504b269ed6d3ec902de16a012a4f059602be...

Step 7: Manifest Signature

PASS

Manifest hash: bd4a0fa1182c6e84ccc7c9362e5aa05657dcb92fd4b2274f...

Steps: 5 PASS / 1 PARTIAL / 0 FAIL

Ledger hash (signed): c9e522e69faf504b269ed6d3ec902de16a012a4f059602be...

Fundamental Rights Impact Assessment (FRIA)

EU AI Act Article 27 | Agent: GovernmentComplianceAgent v1.1 | 2026-04-19

EU AI Act Risk Category: STANDARD | Not classified as Annex III high-risk by default

General-purpose AI deployment -- deployer must classify

Fundamental Rights Assessment

Human Dignity (Article 1) -- impact: LOW

- Basis: Advisory output only -- no autonomous enforcement
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + Injection detection active in SECURITY rules

Privacy and Data (Article 8) -- impact: LOW

- Basis: Session data held in memory, not persisted by default
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + Injection detection active in SECURITY rules

Non-Discrimination (Article 21) -- impact: LOW

- Basis: Rate limits apply equally to all callers -- no exemptions
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + Injection detection active in SECURITY rules

Freedom of Expression (Article 11) -- impact: LOW

- Basis: Model cannot be silenced by operator prompt injection
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + Injection detection active in SECURITY rules

Right to Remedy (Article 47) -- impact: LOW

- Basis: Human review gate (HUMAN_REVIEW) available for all decisions
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + Injection detection active in SECURITY rules

Presumption of Innocence (Article 48) -- impact: LOW

- Basis: Agent cannot issue binding sanctions -- advisory role only
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + Injection detection active in SECURITY rules

Technical Mitigations (auto-populated from certification)

Constitutional constraints (CANNOT_MUTATE): agent, version, goal, security, trust_level

Security rules: 23 | Runtime layers: 5

HUMAN_REVIEW triggers: 9 | Timeout: 24h | Block: True

DoS rate limiting: True | Supply chain integrity: True

DEPLOYER ATTESTATION REQUIRED -- see {agent}_fria_{ts}.json for full template

organisation: PLACEHOLDER -- deployer organisation name

completed_by: PLACEHOLDER -- name and role

date: PLACEHOLDER -- ISO 8601 date

review_period: PLACEHOLDER -- e.g. annual review or on material change

signature: PLACEHOLDER -- wet or electronic signature

FRIA template auto-generated -- deployer must complete PLACEHOLDER fields and obtain sign-off before deployment in regulated context.