

AXIOM Agent Certification Report

Generated: 2026-04-18T22:48:48.927132+00:00 | axiom-lang v1.8.0

Agent: Rewriter v1.2

Conformance Level: NON-CONFORMANT

Honesty Rate: 100% (integrity check -- independent of benchmark)

Certification Steps

Step 1: Structural Validation	FAIL
Validator: valid Errors: 0 Warnings: 0 Supply chain: TAMPERED SHA-256: cf65cf1fc91890d5a03450b36fed6900...	
Step 2: Security Stack Audit	PASS
Security score: 4/7 Rules: 5 Runtime layers active: 5/5 [-] HighRiskInput WHEN activation [-] Sandbox DELEGATES route [+] CANNOT_MUTATE declared [+] TRUST_LEVEL declared [-] SANDBOX_AGENT declared	
Step 3: Benchmark Evidence	PASS
v1_4_when_delegates: 93/94 (99%) -- v1_4_when_delegates_20260418_151200.json	
Step 4: Constitutional Integrity	PASS
Protected fields: agent, goal, version, security, trust_level SHA-256: cf65cf1fc91890d5a03450b36fed6900e5b6e6d494cb4b44...	
Step 5: Audit Trail	PARTIAL
History file: rewriter_history.jsonl Entries: 0 Latest:	
Step 6: Honesty Integrity	PASS
Latest run: 100% (40 evals) H:40 S:0 D:0 Ledger overall: 44% (200 total evals) Ledger hash: 2c8b718694046aa4f8559709ab5c1e7d227f519c0095af99...	
Step 7: Manifest Signature	PASS
Manifest hash: 2e65e877721f4ddfc0eccc92fdf368c09b89fea1f6f91bbe... Steps: 4 PASS / 1 PARTIAL / 1 FAIL Ledger hash (signed): 2c8b718694046aa4f8559709ab5c1e7d227f519c0095af99...	

Fundamental Rights Impact Assessment (FRIA)

EU AI Act Article 27 | Agent: Rewriter v1.2 | 2026-04-18

EU AI Act Risk Category: STANDARD | Not classified as Annex III high-risk by default

General-purpose AI deployment -- deployer must classify

Fundamental Rights Assessment

Human Dignity (Article 1) -- impact: LOW

- Basis: Advisory output only -- no autonomous enforcement
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + Injection detection active in SECURITY rules

Privacy and Data (Article 8) -- impact: LOW

- Basis: Session data held in memory, not persisted by default
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + Injection detection active in SECURITY rules

Non-Discrimination (Article 21) -- impact: LOW

- Basis: Rate limits apply equally to all callers -- no exemptions
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + Injection detection active in SECURITY rules

Freedom of Expression (Article 11) -- impact: LOW

- Basis: Model cannot be silenced by operator prompt injection
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + Injection detection active in SECURITY rules

Right to Remedy (Article 47) -- impact: LOW

- Basis: Human review gate (HUMAN_REVIEW) available for all decisions
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + Injection detection active in SECURITY rules

Presumption of Innocence (Article 48) -- impact: LOW

- Basis: Agent cannot issue binding sanctions -- advisory role only
- + Security block is constitutionally protected (CANNOT_MUTATE)
 - + Injection detection active in SECURITY rules

Technical Mitigations (auto-populated from certification)

Constitutional constraints (CANNOT_MUTATE): agent, goal, version, security, trust_level

Security rules: 5 | Runtime layers: 5

HUMAN_REVIEW triggers: 8 | Timeout: 24h | Block: True

DoS rate limiting: True | Supply chain integrity: True

DEPLOYER ATTESTATION REQUIRED -- see {agent}_fria_{ts}.json for full template

organisation: PLACEHOLDER -- deployer organisation name

completed_by: PLACEHOLDER -- name and role

date: PLACEHOLDER -- ISO 8601 date

review_period: PLACEHOLDER -- e.g. annual review or on material change

signature: PLACEHOLDER -- wet or electronic signature

FRIA template auto-generated -- deployer must complete PLACEHOLDER fields and obtain sign-off before deployment in regulated context.