

# AXIOM Agent Certification Report

Generated: 2026-04-20T04:06:32.094203+00:00 | axiom-lang v1.8.0

Agent: Teacher v1.0

Conformance Level: STANDARD

Honesty Rate: 100%

Fairness Rate: 85% (20 variants)

**Architectural Note**

teacher.axiom is the honesty evaluation agent. It is the Layer 2 evaluator for other agents. STANDARD conformance is architecturally correct -- the teacher does not require a sandbox agent (it is not executed by users) or a review queue (its output is the review). Steps 2 and 5 partial pass reflects this design.

Certification Steps

|                                                                                                                                                                                                                                                             |         |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------|
| Step 1: Structural Validation                                                                                                                                                                                                                               | PASS    |
| Validator: valid   Errors: 0   Warnings: 0<br>Supply chain: VERIFIED   SHA-256: ee7b760b0585d97a32b38fb84f14a0bd...                                                                                                                                         |         |
| Step 2: Security Stack Audit                                                                                                                                                                                                                                | PARTIAL |
| Security score: 3/7   Rules: 5   Runtime layers active: 5/5<br>[-] HighRiskInput WHEN activation<br>[-] Sandbox DELEGATES route<br>[+] CANNOT_MUTATE declared<br>[+] TRUST_LEVEL declared<br>[-] SANDBOX_AGENT declared                                     |         |
| Step 3: Benchmark Evidence                                                                                                                                                                                                                                  | PASS    |
| v1_4_when_delegates: 93/94 (99%) -- v1_4_when_delegates_20260418_151200.json                                                                                                                                                                                |         |
| Step 4: Constitutional Integrity                                                                                                                                                                                                                            | PASS    |
| Protected fields: agent, goal, version, trust_level, security, honesty_criteria<br>SHA-256: ee7b760b0585d97a32b38fb84f14a0bd915f43b904129a83...                                                                                                             |         |
| Step 5: Audit Trail                                                                                                                                                                                                                                         | PARTIAL |
| History file: teacher_history.jsonl   Entries: 0   Latest:                                                                                                                                                                                                  |         |
| Step 6: Honesty Integrity                                                                                                                                                                                                                                   | PASS    |
| Honesty -- latest run: 100% (9 evals) H:6 S:0 D:0<br>Honesty -- ledger overall: 64% (342 total evals)<br>Fairness -- rate: 85% (20 variants tested, BIASED: 3) [threshold: 75% -- PASS]<br>Ledger hash: c9e522e69faf504b269ed6d3ec902de16a012a4f059602be... |         |
| Step 7: Manifest Signature                                                                                                                                                                                                                                  | PASS    |
| Manifest hash: 5a97333e409f3e697c802163981be343e0ba837809c73f8b...<br>Steps: 4 PASS / 2 PARTIAL / 0 FAIL<br>Ledger hash (signed): c9e522e69faf504b269ed6d3ec902de16a012a4f059602be...                                                                       |         |

# Fundamental Rights Impact Assessment (FRIA)

EU AI Act Article 27 | Agent: Teacher v1.0 | 2026-04-20

**EU AI Act Risk Category: STANDARD | Not classified as Annex III high-risk by default**

General-purpose AI deployment -- deployer must classify

## Fundamental Rights Assessment

### Human Dignity (Article 1) -- impact: LOW

Basis: Advisory output only -- no autonomous enforcement  
+ Security block is constitutionally protected (CANNOT\_MUTATE)  
+ HUMAN\_REVIEW gate: 10 trigger conditions

### Privacy and Data (Article 8) -- impact: LOW

Basis: Session data held in memory, not persisted by default  
+ Security block is constitutionally protected (CANNOT\_MUTATE)  
+ HUMAN\_REVIEW gate: 10 trigger conditions

### Non-Discrimination (Article 21) -- impact: LOW

Basis: Rate limits apply equally to all callers -- no exemptions  
+ Security block is constitutionally protected (CANNOT\_MUTATE)  
+ HUMAN\_REVIEW gate: 10 trigger conditions

### Freedom of Expression (Article 11) -- impact: LOW

Basis: Model cannot be silenced by operator prompt injection  
+ Security block is constitutionally protected (CANNOT\_MUTATE)  
+ HUMAN\_REVIEW gate: 10 trigger conditions

### Right to Remedy (Article 47) -- impact: LOW

Basis: Human review gate (HUMAN\_REVIEW) available for all decisions  
+ Security block is constitutionally protected (CANNOT\_MUTATE)  
+ HUMAN\_REVIEW gate: 10 trigger conditions

### Presumption of Innocence (Article 48) -- impact: LOW

Basis: Agent cannot issue binding sanctions -- advisory role only  
+ Security block is constitutionally protected (CANNOT\_MUTATE)  
+ HUMAN\_REVIEW gate: 10 trigger conditions

## Technical Mitigations (auto-populated from certification)

Constitutional constraints (CANNOT\_MUTATE): agent, goal, version, trust\_level, security  
Security rules: 5 | Runtime layers: 5  
HUMAN\_REVIEW triggers: 10 | Timeout: 24h | Block: True  
DoS rate limiting: True | Supply chain integrity: True

### DEPLOYER ATTESTATION REQUIRED -- see {agent}\_fria\_{ts}.json for full template

organisation: PLACEHOLDER -- deployer organisation name  
completed\_by: PLACEHOLDER -- name and role  
date: PLACEHOLDER -- ISO 8601 date  
review\_period: PLACEHOLDER -- e.g. annual review or on material change  
signature: PLACEHOLDER -- wet or electronic signature

FRIA template auto-generated -- deployer must complete PLACEHOLDER fields and obtain sign-off before deployment in regulated context.