



QRF Hypothesis Loop vs Baseline (Arbor-style)



Model: gemma3-1b-srd4 (ollama) Average efficiency: 1.53× † = hard non-canonical task

Canonical algorithms ceiling at 1.0× (baseline already 100%); QRF's advantage appears on hard tasks where single-shot fails. Caveat: QRF spends ~3-5× the tokens (quality, not cost, is plotted).