



QRF Hypothesis Loop vs Baseline (Arbor-style)



Model: gemma3-1b-srd4 (ollama) Median (robust): 1.33× · Mean: 1.53× · 1 rescued † = hard non-canonical task

QRF helps where single-shot leaves headroom; the ratio mean can be inflated by near-zero baselines, so the median is the honest headline. Caveat: QRF spends ~3-5× the tokens (quality, not cost, is plotted).